

# STEP-1 Business Problem Understanding

Predict the Personality types using Machine learning models

## STEP-2 Data Understanding

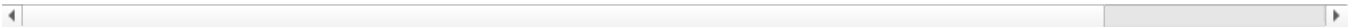
```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.simplefilter("ignore")
```

```
In [2]: df=pd.read_csv("Extrovert vs Introvert.csv")
df
```

```
Out[2]:
```

	Time_spent_Alone	Stage_fear	Social_event_attendance	Going_outside	Drained_after_socializing	Friends_circle_size	Post_fr
0	4.0	No	4.0	6.0	No	13.0	
1	9.0	Yes	0.0	0.0	Yes	0.0	
2	9.0	Yes	1.0	2.0	Yes	5.0	
3	0.0	No	6.0	7.0	No	14.0	
4	3.0	No	9.0	4.0	No	8.0	
...	...	...	...	...	...	...	...
2895	3.0	No	7.0	6.0	No	6.0	
2896	3.0	No	8.0	3.0	No	14.0	
2897	4.0	Yes	1.0	1.0	Yes	4.0	
2898	11.0	Yes	1.0	NaN	Yes	2.0	
2899	3.0	No	6.0	6.0	No	6.0	

2900 rows × 8 columns



```
In [3]: df.shape
```

```
Out[3]: (2900, 8)
```

```
In [4]: df["Time_spent_Alone"]
```

```
Out[4]:
```

0	4.0
1	9.0
2	9.0
3	0.0
4	3.0
...	...
2895	3.0
2896	3.0
2897	4.0
2898	11.0
2899	3.0

Name: Time\_spent\_Alone, Length: 2900, dtype: float64

```
In [5]: df["Time_spent_Alone"].unique()
```

```
Out[5]: array([ 4.,  9.,  0.,  3.,  1.,  2., 10.,  6.,  5.,  8., nan,  7., 11.])
```

```
In [6]: df["Time_spent_Alone"].value_counts()
```

```
Out[6]: Time_spent_Alone
0.0      369
2.0      357
3.0      353
1.0      326
9.0      206
10.0     196
4.0      190
7.0      190
5.0      180
8.0      180
6.0      150
11.0     140
Name: count, dtype: int64
```

```
In [7]: df["Stage_fear"]
```

```
Out[7]: 0      No
1      Yes
2      Yes
3      No
4      No
...
2895   No
2896   No
2897   Yes
2898   Yes
2899   No
Name: Stage_fear, Length: 2900, dtype: object
```

```
In [8]: df["Stage_fear"].unique()
```

```
Out[8]: array(['No', 'Yes', nan], dtype=object)
```

```
In [9]: df["Stage_fear"].value_counts()
```

```
Out[9]: Stage_fear
No      1417
Yes     1410
Name: count, dtype: int64
```

```
In [10]: df["Social_event_attendance"]
```

```
Out[10]: 0      4.0
1      0.0
2      1.0
3      6.0
4      9.0
...
2895   7.0
2896   8.0
2897   1.0
2898   1.0
2899   6.0
Name: Social_event_attendance, Length: 2900, dtype: float64
```

```
In [11]: df["Social_event_attendance"].unique()
```

```
Out[11]: array([ 4.,  0.,  1.,  6.,  9.,  7.,  8.,  3.,  5.,  2., 10., nan])
```

```
In [12]: df["Social_event_attendance"].value_counts()
```

```
Out[12]: Social_event_attendance
2.0      408
0.0      378
1.0      322
3.0      317
4.0      255
6.0      239
7.0      239
9.0      236
5.0      224
8.0      206
10.0      14
Name: count, dtype: int64
```

```
In [13]: df["Going_outside"]
```

```
Out[13]: 0      6.0
         1      0.0
         2      2.0
         3      7.0
         4      4.0
         ...
        2895    6.0
        2896    3.0
        2897    1.0
        2898    NaN
        2899    6.0
        Name: Going_outside, Length: 2900, dtype: float64
```

```
In [14]: df["Going_outside"].unique()
```

```
Out[14]: array([ 6.,  0.,  2.,  7.,  4.,  5., nan,  3.,  1.])
```

```
In [15]: df["Going_outside"].value_counts()
```

```
Out[15]: Going_outside
0.0      498
2.0      456
1.0      429
5.0      374
4.0      359
6.0      335
3.0      209
7.0      174
        Name: count, dtype: int64
```

```
In [16]: df["Drained_after_socializing"]
```

```
Out[16]: 0      No
         1      Yes
         2      Yes
         3      No
         4      No
         ...
        2895    No
        2896    No
        2897    Yes
        2898    Yes
        2899    No
        Name: Drained_after_socializing, Length: 2900, dtype: object
```

```
In [17]: df["Drained_after_socializing"].unique()
```

```
Out[17]: array(['No', 'Yes', nan], dtype=object)
```

```
In [18]: df["Drained_after_socializing"].value_counts()
```

```
Out[18]: Drained_after_socializing
No      1441
Yes     1407
        Name: count, dtype: int64
```

```
In [19]: df["Friends_circle_size"]
```

```
Out[19]: 0      13.0
         1      0.0
         2      5.0
         3     14.0
         4      8.0
         ...
        2895    6.0
        2896   14.0
        2897    4.0
        2898    2.0
        2899    6.0
        Name: Friends_circle_size, Length: 2900, dtype: float64
```

```
In [20]: df["Friends_circle_size"].unique()
```

```
Out[20]: array([13.,  0.,  5., 14.,  8.,  6.,  7., 15.,  4., 10.,  1., 12.,  2.,
        11.,  9.,  3., nan])
```

```
In [21]: df["Friends_circle_size"].value_counts()
```

```
Out[21]: Friends_circle_size
5.0      301
3.0      283
1.0      281
2.0      274
4.0      254
8.0      165
12.0     148
10.0     146
14.0     144
6.0      137
9.0      135
11.0     134
7.0      133
13.0     123
0.0      106
15.0      59
Name: count, dtype: int64
```

```
In [22]: df["Post_frequency"]
```

```
Out[22]: 0      5.0
1      3.0
2      2.0
3      8.0
4      5.0
...
2895   6.0
2896   9.0
2897   0.0
2898   0.0
2899   9.0
Name: Post_frequency, Length: 2900, dtype: float64
```

```
In [23]: df["Post_frequency"].unique()
```

```
Out[23]: array([ 5.,  3.,  2.,  8.,  6.,  7.,  0., 10.,  4.,  1.,  9., nan])
```

```
In [24]: df["Post_frequency"].value_counts()
```

```
Out[24]: Post_frequency
2.0      481
1.0      455
0.0      451
7.0      236
5.0      212
6.0      210
3.0      208
4.0      195
8.0      193
9.0      171
10.0      23
Name: count, dtype: int64
```

```
In [25]: df["Personality"]
```

```
Out[25]: 0      Extrovert
1      Introvert
2      Introvert
3      Extrovert
4      Extrovert
...
2895   Extrovert
2896   Extrovert
2897   Introvert
2898   Introvert
2899   Extrovert
Name: Personality, Length: 2900, dtype: object
```

```
In [26]: df["Personality"].unique()
```

```
Out[26]: array(['Extrovert', 'Introvert'], dtype=object)
```

```
In [27]: df["Personality"].value_counts()
```

```
Out[27]: Personality
Extrovert    1491
Introvert    1409
Name: count, dtype: int64
```

```
In [28]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2900 entries, 0 to 2899
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Time_spent_Alone                      2837 non-null   float64
1   Stage_fear                           2827 non-null   object
2   Social_event_attendance               2838 non-null   float64
3   Going_outside                        2834 non-null   float64
4   Drained_after_socializing             2848 non-null   object
5   Friends_circle_size                  2823 non-null   float64
6   Post_frequency                       2835 non-null   float64
7   Personality                           2900 non-null   object
dtypes: float64(5), object(3)
memory usage: 181.4+ KB

```

```

In [29]: continuous = ["Time_spent_Alone", "Social_event_attendance", "Going_outside", "Friends_circle_size", "Post_frequency"]
discrete_categorical = ["Stage_fear", "Drained_after_socializing", "Personality"]

```

## Exploratory Data Analysis

### For Continuous Variables

```

In [30]: df[continuous].describe()

```

```

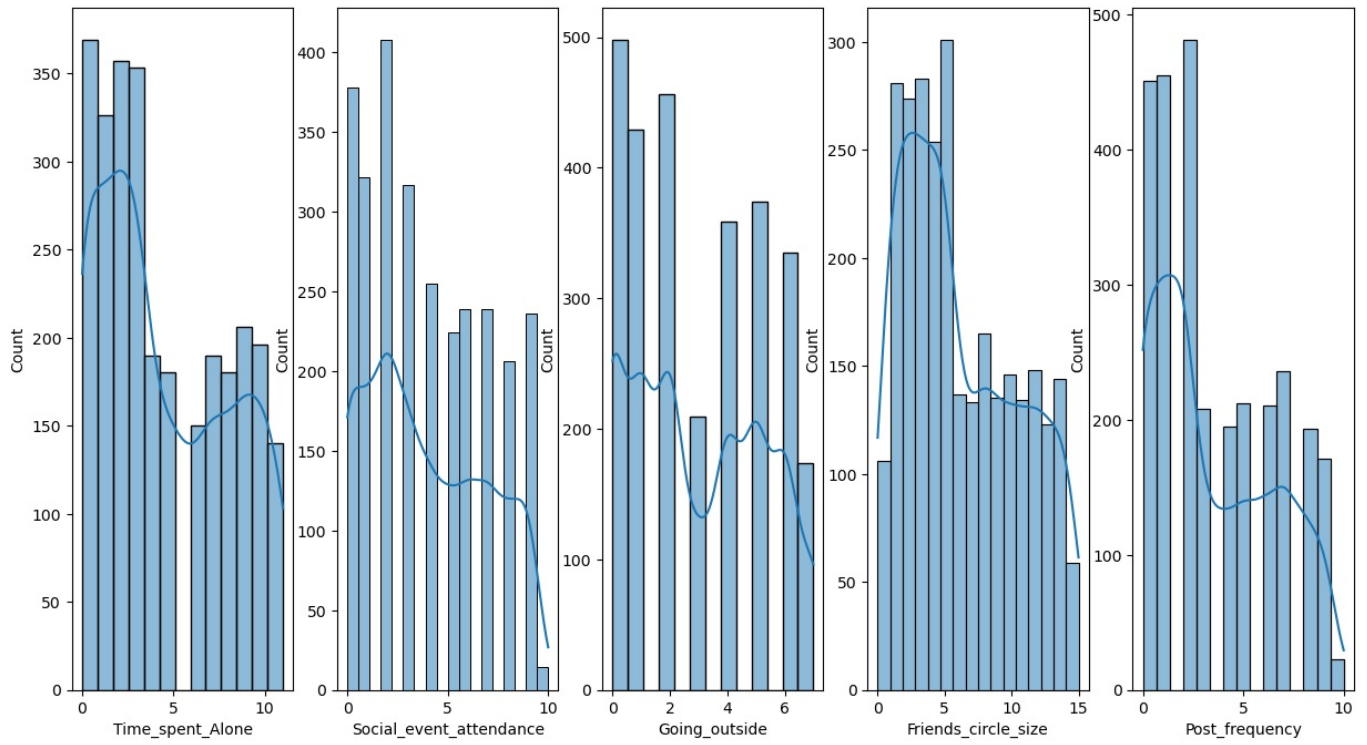
Out[30]:
```

	Time_spent_Alone	Social_event_attendance	Going_outside	Friends_circle_size	Post_frequency
count	2837.000000	2838.000000	2834.000000	2823.000000	2835.000000
mean	4.505816	3.963354	3.000000	6.268863	3.564727
std	3.479192	2.903827	2.247327	4.289693	2.926582
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	2.000000	2.000000	1.000000	3.000000	1.000000
50%	4.000000	3.000000	3.000000	5.000000	3.000000
75%	8.000000	6.000000	5.000000	10.000000	6.000000
max	11.000000	10.000000	7.000000	15.000000	10.000000

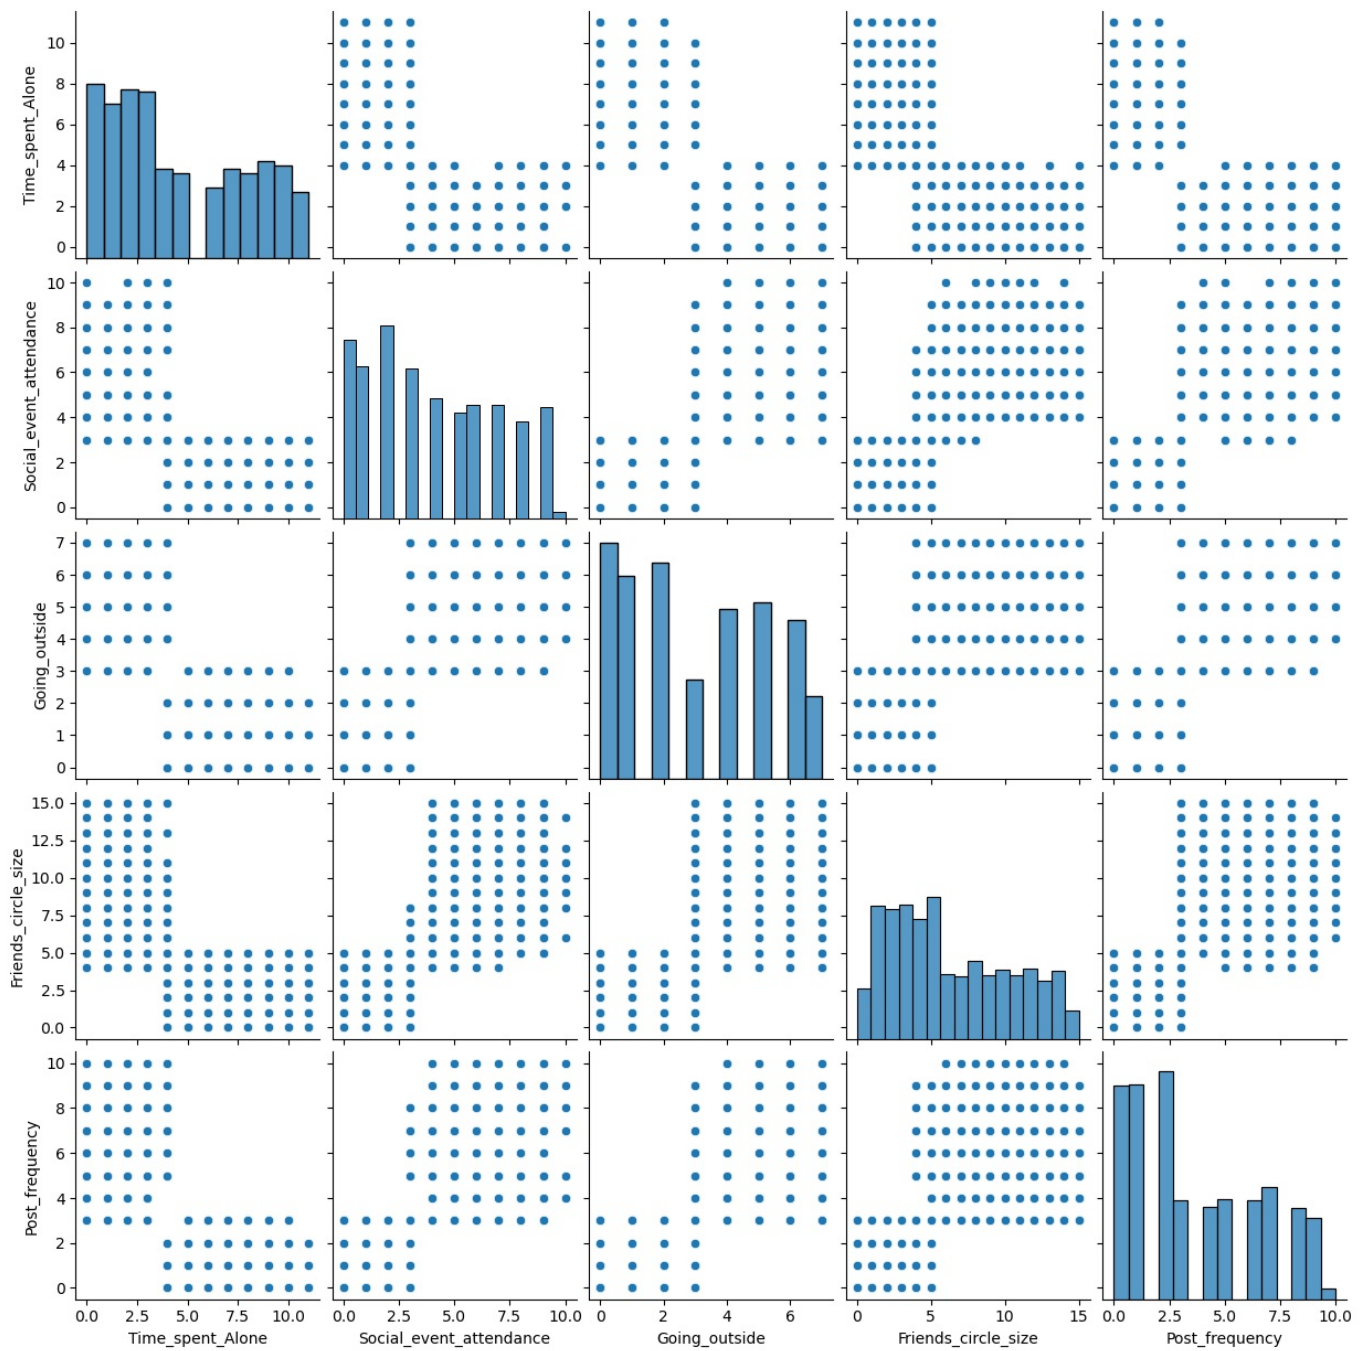
```

In [31]: plt.rcParams["figure.figsize"] = (18,8)
plt.subplot(1,6, 1)
sns.histplot(df["Time_spent_Alone"],kde=True)
plt.subplot(1,6, 2)
sns.histplot(df["Social_event_attendance"],kde=True)
plt.subplot(1,6, 3)
sns.histplot(df["Going_outside"],kde=True)
plt.subplot(1,6, 4)
sns.histplot(df["Friends_circle_size"],kde=True)
plt.subplot(1,6, 5)
sns.histplot(df["Post_frequency"],kde=True)
plt.suptitle("Univariate Analysis on Numerical Columns")
plt.show()

```



```
In [32]: sns.pairplot(df[continuous])  
plt.show()
```




```
In [33]: df[continuous].corr()
```

```
Out[33]:
```

	Time_spent_Alone	Social_event_attendance	Going_outside	Friends_circle_size	Post_frequency
Time_spent_Alone	1.000000	-0.733011	-0.750760	-0.717185	-0.732649
Social_event_attendance	-0.733011	1.000000	0.747756	0.734795	0.744615
Going_outside	-0.750760	0.747756	1.000000	0.736390	0.770819
Friends_circle_size	-0.717185	0.734795	0.736390	1.000000	0.707888
Post_frequency	-0.732649	0.744615	0.770819	0.707888	1.000000

```
In [34]: sns.heatmap(df[continuous].corr(),annot=True)
plt.show()
```



## For Discrete variables

```
In [35]: df[discrete_categorical].describe()
```

```
Out[35]:
```

	Stage_fear	Drained_after_socializing	Personality
count	2827	2848	2900
unique	2	2	2
top	No	No	Extrovert
freq	1417	1441	1491

## Steps to be followed for data cleaning

### Check for Wrong Data

### Check for Wrong Data type

```
In [36]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2900 entries, 0 to 2899
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Time_spent_Alone      2837 non-null   float64
1   Stage_fear            2827 non-null   object
2   Social_event_attendance  2838 non-null   float64
3   Going_outside         2834 non-null   float64
4   Drained_after_socializing  2848 non-null   object
5   Friends_circle_size    2823 non-null   float64
6   Post_frequency        2835 non-null   float64
7   Personality           2900 non-null   object
dtypes: float64(5), object(3)
memory usage: 181.4+ KB
```

### Check for duplicates

```
In [37]: df.duplicated().sum()
```

```
Out[37]: 388
```

```
In [38]: df[df.duplicated()]
```



Out[38]:

	Time_spent_Alone	Stage_fear	Social_event_attendance	Going_outside	Drained_after_socializing	Friends_circle_size	Post_fr
47	10.0	Yes	1.0	2.0	Yes	2.0	
217	5.0	Yes	2.0	0.0	Yes	2.0	
246	9.0	Yes	0.0	1.0	Yes	2.0	
248	9.0	Yes	0.0	2.0	Yes	3.0	
254	7.0	Yes	0.0	0.0	Yes	3.0	
...	...	...	...	...	...	...	...
2884	11.0	Yes	0.0	2.0	Yes	3.0	
2890	8.0	Yes	2.0	0.0	Yes	1.0	
2891	6.0	Yes	3.0	1.0	Yes	5.0	
2892	9.0	Yes	2.0	0.0	Yes	1.0	
2895	3.0	No	7.0	6.0	No	6.0	

388 rows × 8 columns



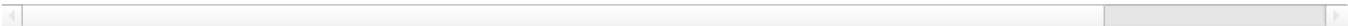
In [39]:

```
df[~df.duplicated()]
```

Out[39]:

	Time_spent_Alone	Stage_fear	Social_event_attendance	Going_outside	Drained_after_socializing	Friends_circle_size	Post_fr
0	4.0	No	4.0	6.0	No	13.0	
1	9.0	Yes	0.0	0.0	Yes	0.0	
2	9.0	Yes	1.0	2.0	Yes	5.0	
3	0.0	No	6.0	7.0	No	14.0	
4	3.0	No	9.0	4.0	No	8.0	
...	...	...	...	...	...	...	...
2894	0.0	No	9.0	3.0	No	12.0	
2896	3.0	No	8.0	3.0	No	14.0	
2897	4.0	Yes	1.0	1.0	Yes	4.0	
2898	11.0	Yes	1.0	NaN	Yes	2.0	
2899	3.0	No	6.0	6.0	No	6.0	

2512 rows × 8 columns



### Check for Missing values

In [40]:

```
df.isnull().sum()
```

Out[40]:

Time_spent_Alone	63
Stage_fear	73
Social_event_attendance	62
Going_outside	66
Drained_after_socializing	52
Friends_circle_size	77
Post_frequency	65
Personality	0
dtype:	int64

In [41]:

```
df.isnull().sum()/len(df)*100
```

Out[41]:

Time_spent_Alone	2.172414
Stage_fear	2.517241
Social_event_attendance	2.137931
Going_outside	2.275862
Drained_after_socializing	1.793103
Friends_circle_size	2.655172
Post_frequency	2.241379
Personality	0.000000
dtype:	float64

### Check for skewness

In [42]:

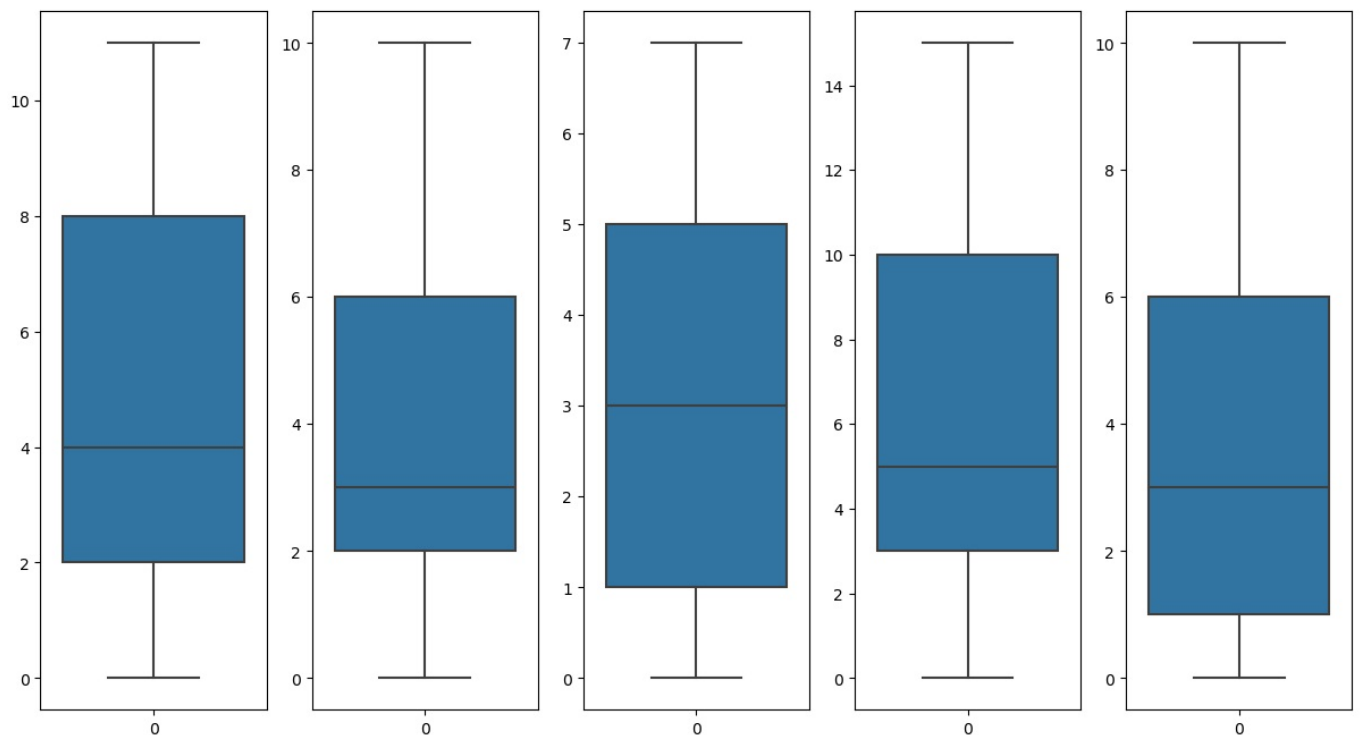
```
df[["Time_spent_Alone","Social_event_attendance","Going_outside","Friends_circle_size","Post_frequency"]].skew()
```

```
Out[42]: Time_spent_Alone      0.385821
Social_event_attendance      0.294742
Going_outside                 0.192891
Friends_circle_size           0.425051
Post_frequency                 0.474510
dtype: float64
```

## Check for outliers

```
In [43]: # Lets visualize the outliers using Boxplot
plt.subplot(1,6, 1)
sns.boxplot(df["Time_spent_Alone"])
plt.subplot(1,6, 2)
sns.boxplot(df["Social_event_attendance"])
plt.subplot(1,6, 3)
sns.boxplot(df["Going_outside"])
plt.subplot(1,6, 4)
sns.boxplot(df["Friends_circle_size"])
plt.subplot(1,6, 5)
sns.boxplot(df["Post_frequency"])
plt.suptitle("Outliers in the data")
plt.show()
```

Outliers in the data



## STEP-3 Data Preprocessing

### I.Data Cleaning

No Treatment for wrong data

No Treatment for wrong datatype

Treating Duplicates

```
In [44]: df.drop_duplicates(inplace=True,ignore_index=True)
df
```

Out [44]:

	Time_spent_Alone	Stage_fear	Social_event_attendance	Going_outside	Drained_after_socializing	Friends_circle_size	Post_fr
0	4.0	No	4.0	6.0	No	13.0	
1	9.0	Yes	0.0	0.0	Yes	0.0	
2	9.0	Yes	1.0	2.0	Yes	5.0	
3	0.0	No	6.0	7.0	No	14.0	
4	3.0	No	9.0	4.0	No	8.0	
...	...	...	...	...	...	...	...
2507	0.0	No	9.0	3.0	No	12.0	
2508	3.0	No	8.0	3.0	No	14.0	
2509	4.0	Yes	1.0	1.0	Yes	4.0	
2510	11.0	Yes	1.0	NaN	Yes	2.0	
2511	3.0	No	6.0	6.0	No	6.0	

2512 rows × 8 columns

Treating Missing values

In [45]:

df.dropna(inplace=True)  
df

Out [45]:

	Time_spent_Alone	Stage_fear	Social_event_attendance	Going_outside	Drained_after_socializing	Friends_circle_size	Post_fr
0	4.0	No	4.0	6.0	No	13.0	
1	9.0	Yes	0.0	0.0	Yes	0.0	
2	9.0	Yes	1.0	2.0	Yes	5.0	
3	0.0	No	6.0	7.0	No	14.0	
4	3.0	No	9.0	4.0	No	8.0	
...	...	...	...	...	...	...	...
2504	5.0	Yes	0.0	1.0	Yes	1.0	
2505	6.0	Yes	0.0	0.0	Yes	3.0	
2508	3.0	No	8.0	3.0	No	14.0	
2509	4.0	Yes	1.0	1.0	Yes	4.0	
2511	3.0	No	6.0	6.0	No	6.0	

2098 rows × 8 columns

In [46]:

df.isnull().sum()

Out [46]:

Time_spent_Alone	0
Stage_fear	0
Social_event_attendance	0
Going_outside	0
Drained_after_socializing	0
Friends_circle_size	0
Post_frequency	0
Personality	0
dtype:	int64

No Treatment for Outliers

II.Data Wrangling

Feature Encoding

In [47]:

df["Stage\_fear"] = df["Stage\_fear"].replace({"No":0,"Yes":1})  
df["Drained\_after\_socializing"] = df["Drained\_after\_socializing"].replace({"No":0,"Yes":1})  
df["Personality"] = df["Personality"].replace({"Extrovert":0,"Introvert":1})

No Feature Transformations

## No Feature Scaling

In [48]:

```
df
```

Out[48]:

	Time_spent_Alone	Stage_fear	Social_event_attendance	Going_outside	Drained_after_socializing	Friends_circle_size	Post_fr
0	4.0	0	4.0	6.0	0	13.0	
1	9.0	1	0.0	0.0	1	0.0	
2	9.0	1	1.0	2.0	1	5.0	
3	0.0	0	6.0	7.0	0	14.0	
4	3.0	0	9.0	4.0	0	8.0	
...	...	...	...	...	...	...	...
2504	5.0	1	0.0	1.0	1	1.0	
2505	6.0	1	0.0	0.0	1	3.0	
2508	3.0	0	8.0	3.0	0	14.0	
2509	4.0	1	1.0	1.0	1	4.0	
2511	3.0	0	6.0	6.0	0	6.0	

2098 rows × 8 columns

## X&y

In [49]:

```
X = df.drop("Personality",axis=1)
y = df["Personality"]
```

## Identify the best random number

In [50]:

```
Train = []
Test = []
CV = []
for i in range(0,101):
    from sklearn.model_selection import train_test_split
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=i)

    from sklearn.linear_model import LogisticRegression
    log_default = LogisticRegression()
    log_default.fit(X_train,y_train)

    ypred_train = log_default.predict(X_train)
    ypred_test = log_default.predict(X_test)
    from sklearn.metrics import accuracy_score
    Train.append(accuracy_score(y_train, ypred_train))
    Test.append(accuracy_score(y_test, ypred_test))

    from sklearn.model_selection import cross_val_score
    CV.append(cross_val_score(log_default, X_train, y_train, cv=5, scoring="accuracy").mean())
#Storing all results
em = pd.DataFrame({"Train":Train, "Test":Test, "CV":CV})
#Find the best random state
gm = em[(abs(em["Train"]-em["Test"])<=0.05) & (abs(em["Test"]-em["CV"])<=0.05)]
#pick the highest CV
rs = gm[gm["CV"]==gm["CV"].max()].index.to_list()[0]
print("best random_state number:",rs)
```

best random\_state number: 57

## III.train\_test\_split

In [51]:

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=57)
```

## Step-4 ML Modelling

In [52]:

```
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier
```

```
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score
from sklearn.model_selection import cross_val_score
```

## 1.Logistic Regression Algorithm

```
In [53]: #Modelling
log_model = LogisticRegression()
log_model.fit(X_train,y_train)

#Evaluation
ypred_train = log_model.predict(X_train)
ypred_test = log_model.predict(X_test)

print("Train Accuracy :",accuracy_score(y_train,ypred_train))
print("cross validation score:",cross_val_score(log_model,X_train,y_train,cv=5,scoring="accuracy").mean())
print("Test Accuracy :",accuracy_score(y_test,ypred_test))
```

Train Accuracy : 0.9249106078665077  
cross validation score: 0.9249129353233831  
Test Accuracy : 0.8761904761904762

## 2.KNN Classifier Algorithm

```
In [54]: #Hyper Parameter Tuning
estimator = KNeighborsClassifier()
param_grid = {"n_neighbors": list(range(1,50))}
knn_grid = GridSearchCV(estimator, param_grid,scoring="accuracy",cv=5)
knn_grid.fit(X_train,y_train)
knn_model = knn_grid.best_estimator_
knn_model
```

```
Out[54]: ▼ KNeighborsClassifier
KNeighborsClassifier(n_neighbors=13)
```

```
In [55]: #Modelling
knn_model = KNeighborsClassifier(n_neighbors=13)
knn_model.fit(X_train,y_train)

#Evaluation
ypred_train = knn_model.predict(X_train)
ypred_test = knn_model.predict(X_test)

print("Train Accuracy :",accuracy_score(y_train,ypred_train))
print("cross validation score :",cross_val_score(knn_model,X_train,y_train,cv=5,scoring="accuracy").mean())
print("Test Accuracy :",accuracy_score(y_test,ypred_test))
```

Train Accuracy : 0.9249106078665077  
cross validation score : 0.9249129353233831  
Test Accuracy : 0.8761904761904762

## 3.Support Vector Machine Algorithm

```
In [56]: #Hyper Parameter Tuning
estimator = SVC()
param_grid = {"C":[0.01,0.1,1],"kernel":["linear","rbf","sigmoid","poly"]}
svm_grid = GridSearchCV(estimator,param_grid,scoring="accuracy",cv=5)
svm_grid.fit(X_train,y_train)
svm_model = svm_grid.best_estimator_
svm_model
```

```
Out[56]: ▼ SVC
SVC(C=0.01, kernel='linear')
```

```
In [57]: #Modelling
svm_model = SVC(C=0.01, kernel="linear")
svm_model.fit(X_train,y_train)

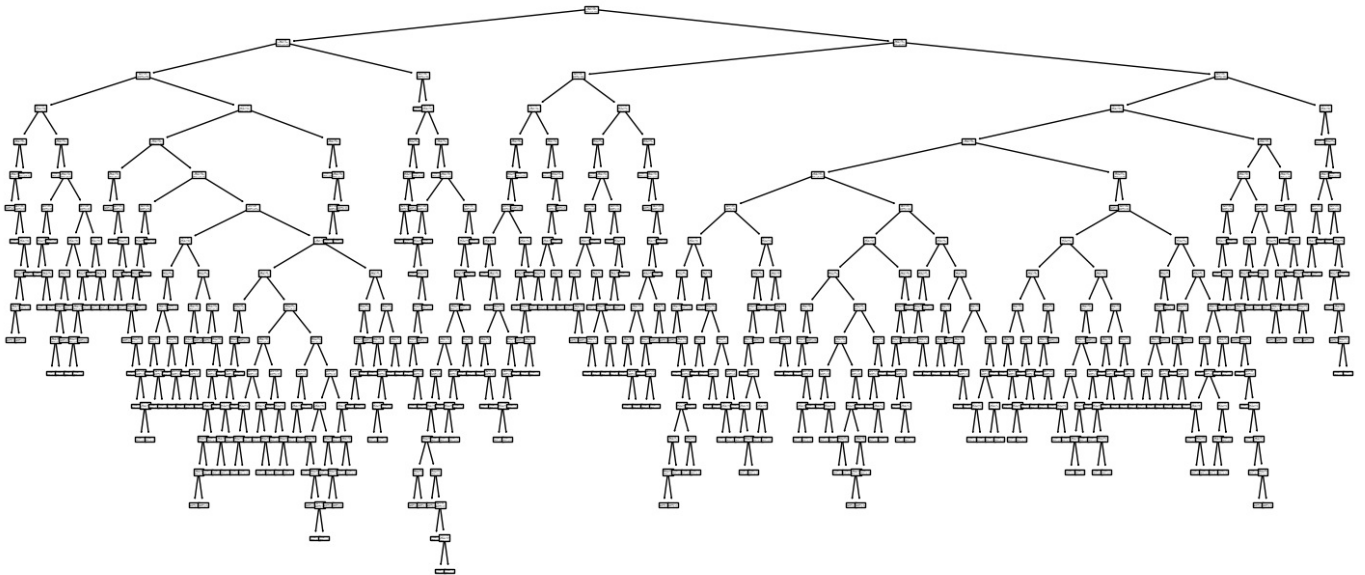
#Evaluation
ypred_train = svm_model.predict(X_train)
ypred_test = svm_model.predict(X_test)
```

```
print("Train Accuracy :",accuracy_score(y_train,ypred_train))
print("cross validation score :",cross_val_score(svm_model,X_train,y_train,cv=5,scoring="accuracy").mean())
print("Test Accuracy :",accuracy_score(y_test,ypred_test))
```

Train Accuracy : 0.9249106078665077  
cross validation score : 0.9249129353233831  
Test Accuracy : 0.8761904761904762

## 4.Decision Tree Classifier Algorithm

```
In [58]: model = DecisionTreeClassifier(random_state=True)
model.fit(X_train,y_train)
from sklearn.tree import plot_tree
plot_tree(model)
plt.show()
```



```
In [59]: #Hyper Parameter Tuning
estimator = DecisionTreeClassifier(random_state=True)
param_grid = {"criterion":["gini", "entropy"],
              "max_depth":list(range(1,16))}
dt_grid = GridSearchCV(estimator,param_grid,scoring="accuracy",cv=5)
dt_grid.fit(X_train,y_train)
dt = dt_grid.best_estimator_
dt
```

```
Out[59]: ▾ DecisionTreeClassifier
DecisionTreeClassifier(max_depth=1, random_state=True)
```

```
In [60]: dt.feature_importances_
```

```
Out[60]: array([0., 0., 0., 0., 1., 0., 0.])
```

```
In [61]: #Important features
feats_dt = pd.DataFrame(data=dt.feature_importances_,
                        index=X.columns,
                        columns=["Importance"])
important_features_dt = feats_dt[feats_dt["Importance"]>0].index.tolist()
important_features_dt
```

```
Out[61]: ['Drained_after_socializing']
```

## Creating Decision Tree model with important parameters and important features

```
In [62]: #Selecting train & test data
X_train_dt = X_train[important_features_dt]
X_test_dt = X_test[important_features_dt]

#Modelling
dt.fit(X_train_dt,y_train)

#Evaluation
ypred_train = dt.predict(X_train_dt)
ypred_test = dt.predict(X_test_dt)

print("Train Accuracy :",accuracy_score(y_train,ypred_train))
```

```
print("cross validation score :",cross_val_score(dt,X_train_dt,y_train,cv=5,scoring="accuracy").mean())
print("Test Accuracy :",accuracy_score(y_test,ypred_test))
```

Train Accuracy : 0.9249106078665077  
cross validation score : 0.9249129353233831  
Test Accuracy : 0.8761904761904762

## 5.Random Forest Classifier Algorithm

```
In [63]: #Hyper Parameter Tuning
estimator = RandomForestClassifier(random_state=True)
param_grid = {"n_estimators":list(range(1,51))}
rf_grid = GridSearchCV(estimator,param_grid,scoring="accuracy",cv=5)
rf_grid.fit(X_train,y_train)
rf = rf_grid.best_estimator_
rf
```

```
Out[63]: ▼ RandomForestClassifier
RandomForestClassifier(n_estimators=43, random_state=True)
```

```
In [64]: #Important features
feats_rf = pd.DataFrame(data=rf.feature_importances_,
                        index=X.columns,
                        columns=["Importance"])
important_features_rf = feats_rf[feats_rf["Importance"]>0].index.tolist()
important_features_rf
```

```
Out[64]: ['Time_spent_Alone',
          'Stage_fear',
          'Social_event_attendance',
          'Going_outside',
          'Drained_after_socializing',
          'Friends_circle_size',
          'Post_frequency']
```

```
In [65]: #Selecting train & test data
X_train_rf = X_train[important_features_rf]
X_test_rf = X_test[important_features_rf]

#Modelling
rf.fit(X_train_rf,y_train)

#Evaluation
ypred_train = rf.predict(X_train_rf)
ypred_test = rf.predict(X_test_rf)

print("Train Accuracy :",accuracy_score(y_train,ypred_train))
print("cross validation score :",cross_val_score(rf,X_train_rf,y_train,cv=5,scoring="accuracy").mean())
print("Test Accuracy :",accuracy_score(y_test,ypred_test))
```

Train Accuracy : 0.9791418355184743  
cross validation score : 0.9010820895522389  
Test Accuracy : 0.8571428571428571

## 6.Ada Boost Classifier Algorithm

```
In [66]: #Hyper Parameter Tuning
estimator = AdaBoostClassifier(random_state=True)
param_grid = {"n_estimators":list(range(1,51))}
ab_grid = GridSearchCV(estimator,param_grid,scoring="accuracy",cv=5)
ab_grid.fit(X_train,y_train)
ab = ab_grid.best_estimator_
ab
```

```
Out[66]: ▼ AdaBoostClassifier
AdaBoostClassifier(n_estimators=1, random_state=True)
```

```
In [67]: #Important features
feats_ab = pd.DataFrame(data=ab.feature_importances_,
                        index=X.columns,
                        columns=["Importance"])
important_features_ab = feats_ab[feats_ab["Importance"]>0].index.tolist()
important_features_ab
```

```
Out[67]: ['Drained_after_socializing']
```

```
In [68]: #Selecting train & test data
```

```
X_train_ab = X_train[important_features_ab]
X_test_ab = X_test[important_features_ab]
```

```
#Modelling
```

```
ab.fit(X_train_ab,y_train)
```

```
#Evaluation
```

```
ypred_train = ab.predict(X_train_ab)
```

```
ypred_test = ab.predict(X_test_ab)
```

```
print("Train Accuracy :",accuracy_score(y_train,ypred_train))
```

```
print("cross validation score :",cross_val_score(ab,X_train_ab,y_train,cv=5,scoring="accuracy").mean())
```

```
print("Test Accuracy :",accuracy_score(y_test,ypred_test))
```

Train Accuracy : 0.9249106078665077

cross validation score : 0.9249129353233831

Test Accuracy : 0.8761904761904762

## 7.Gradient Boosting Classifier Algorithm

In [69]:

```
#Hyper Parameter Tuning
```

```
estimator = GradientBoostingClassifier(random_state=True)
```

```
param_grid = {"n_estimators":list(range(1,10)),
              "learning_rate":[0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1.0]}
```

```
gb_grid = GridSearchCV(estimator,param_grid,scoring="accuracy",cv=5)
```

```
gb_grid.fit(X_train,y_train)
```

```
gb = gb_grid.best_estimator_
gb
```

Out[69]:

```
▼ GradientBoostingClassifier
GradientBoostingClassifier(learning_rate=0.8, n_estimators=8, random_state=True)
```

In [70]:

```
#Important features
```

```
feats_gb = pd.DataFrame(data=gb.feature_importances_,
                        index=X.columns,
                        columns=["Importance"])
```

```
important_features_gb = feats_gb[feats_gb["Importance"]>0].index.tolist()
important_features_gb
```

Out[70]:

```
['Time_spent_Alone',
 'Stage_fear',
 'Social_event_attendance',
 'Going_outside',
 'Drained_after_socializing',
 'Friends_circle_size',
 'Post_frequency']
```

In [71]:

```
#Selecting train & test data
```

```
X_train_gb = X_train[important_features_gb]
```

```
X_test_gb = X_test[important_features_gb]
```

```
#Modelling
```

```
gb.fit(X_train_gb,y_train)
```

```
#Evaluation
```

```
ypred_train = gb.predict(X_train_gb)
```

```
ypred_test = gb.predict(X_test_gb)
```

```
print("Train Accuracy :",accuracy_score(y_train,ypred_train))
```

```
print("cross validation score :",cross_val_score(gb,X_train_gb,y_train,cv=5,scoring="accuracy").mean())
```

```
print("Test Accuracy :",accuracy_score(y_test,ypred_test))
```

Train Accuracy : 0.9249106078665077

cross validation score : 0.9255099502487563

Test Accuracy : 0.8761904761904762

## 8.XG Boost Classifier Algorithm

In [72]:

```
#Hyper Parameter Tuning
```

```
estimator = XGBClassifier()
```

```
param_grid = {"n_estimators":[10,20,40,100],
              "max_depth":[3,4,5],
              "gamma":[0,0.15,0.3,0.5,1]}
```

```
xgb_grid = GridSearchCV(estimator,param_grid,scoring="accuracy",cv=5)
```

```
xgb_grid.fit(X_train,y_train)
```

```
xgb = xgb_grid.best_estimator_
```



xgb

```
Out[72]: XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=0, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
```

```
In [73]: #Important features
feats_xgb = pd.DataFrame(data=xgb.feature_importances_,
                        index=X.columns,
                        columns=["Importance"])
important_features_xgb = feats_xgb[feats_xgb["Importance"]>0].index.tolist()
important_features_xgb
```

```
Out[73]: ['Time_spent_Alone',
          'Stage_fear',
          'Social_event_attendance',
          'Going_outside',
          'Friends_circle_size',
          'Post_frequency']
```

```
In [74]: #Selecting train & test data
X_train_xgb = X_train[important_features_xgb]
X_test_xgb = X_test[important_features_xgb]

#Modelling
xgb.fit(X_train_xgb,y_train)

#Evaluation
ypred_train = xgb.predict(X_train_xgb)
ypred_test = xgb.predict(X_test_xgb)

print("Train Accuracy :",accuracy_score(y_train,ypred_train))
print("cross validation score :",cross_val_score(xgb,X_train_xgb,y_train,cv=5,scoring="accuracy").mean())
print("Test Accuracy :",accuracy_score(y_test,ypred_test))
```

Train Accuracy : 0.9249106078665077  
cross validation score : 0.9249129353233831  
Test Accuracy : 0.8761904761904762

## Step-5 Save the Best model

```
In [75]: from joblib import dump
dump(xgb,"Extrovert vs Introvert.joblib")
```

```
Out[75]: ['Extrovert vs Introvert.joblib']
```

## Step-6 Predict on new data

### Person-1:-

```
In [76]: input_data ={"Time_spent_Alone":5,
                    "Stage_fear":"Yes",
                    "Social_event_attendance":0,
                    "Going_outside":2,
                    "Friends_circle_size":1,
                    "Post_frequency":1,
                    }
```

```
In [77]: df = pd.DataFrame(input_data,index=[0])
df
```

```
Out[77]:
```

	Time_spent_Alone	Stage_fear	Social_event_attendance	Going_outside	Friends_circle_size	Post_frequency
0	5	Yes	0	2	1	1

## Apply Data Preprocessing on Unknown data

```
In [78]: df.drop_duplicates(inplace=True,ignore_index=True)
df.dropna(inplace=True)
df["Stage_fear"] = df["Stage_fear"].replace({"No":0,"Yes":1})

X_new = df
```

```
In [79]: X_new
```

```
Out[79]:
```

	Time_spent_Alone	Stage_fear	Social_event_attendance	Going_outside	Friends_circle_size	Post_frequency
0	5	1	0	2	1	1

```
In [80]: X_new = X_new[["Time_spent_Alone","Stage_fear","Social_event_attendance","Going_outside","Friends_circle_size"],
xgb.predict(X_new)
```

```
Out[80]: array([1])
```

This Person is an Introvert

## Person-2:-

```
In [81]: input_data ={"Time_spent_Alone":1,
                    "Stage_fear":"No",
                    "Social_event_attendance":9,
                    "Going_outside":6,
                    "Friends_circle_size":11,
                    "Post_frequency":9,
                    }
```

```
In [82]: df = pd.DataFrame(input_data,index=[0])
df
```

```
Out[82]:
```

	Time_spent_Alone	Stage_fear	Social_event_attendance	Going_outside	Friends_circle_size	Post_frequency
0	1	No	9	6	11	9

## Apply Data Preprocessing on unknown data

```
In [83]: df.drop_duplicates(inplace=True,ignore_index=True)
df.dropna(inplace=True)
df["Stage_fear"] = df["Stage_fear"].replace({"No":0,"Yes":1})

X_new = df
```

```
In [84]: X_new
```

```
Out[84]:
```

	Time_spent_Alone	Stage_fear	Social_event_attendance	Going_outside	Friends_circle_size	Post_frequency
0	1	0	9	6	11	9

```
In [85]: X_new = X_new[["Time_spent_Alone","Stage_fear","Social_event_attendance","Going_outside","Friends_circle_size"],
xgb.predict(X_new)
```

```
Out[85]: array([0])
```

This Person is an Extrovert

```
In [ ]:
```